

Gradient Descent

Dimitris Papadimitriou

INFO 251: Applied Machine Learning

March 4, 2020

Optimization (General Formulation)

- ▶ Most ML problems can be formulated as optimization problems

$$\min_{z_1, z_2, \dots, z_n} J(z_1, z_2, \dots, z_n)$$

- ▶ where $z_1, z_2 \dots z_n$ are our parameters
- ▶ In the regression setting these would be the intercept and the slopes
- ▶ Closed form solution usually not possible →
Gradient Descent

Gradient Descent Intuition

- ▶ Gradient always points in the direction of greatest increase of the function
- ▶ Thus negative gradient points towards steepest descent
- ▶ We expect that taking steps in the direction of the negative derivative will lead us to the minimum
- ▶ How big steps though?

Linear Regression

- ▶ Assume $y_i = \alpha + \beta x_i + \epsilon_i$, $i = 1, \dots, N$
- ▶ Objective of OLS: $J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - \alpha - \beta x_i)^2$
- ▶ Estimate α, β by $\min_{\alpha, \beta} \frac{1}{2N} \sum_{i=1}^N (y_i - \alpha - \beta x_i)^2$
- ▶ "In general we optimize a function with respect to some variables by setting the derivatives w.r.t. those variables equal to zero and solving for those parameters"
- ▶ Closed form solution

$$\begin{bmatrix} \alpha \\ \beta \end{bmatrix} = (X^T X)^{-1} X^T Y$$

where $X = \begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_N \end{bmatrix}$ and $Y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_N \end{bmatrix}$

Issues With That...

- ▶ So why GD if we have closed form solution?

Issues With That...

- ▶ So why GD if we have closed form solution?
- ▶ For OLS inverting $X^T X$ can be very demanding

Issues With That...

- ▶ So why GD if we have closed form solution?
- ▶ For OLS inverting $X^T X$ can be very demanding
- ▶ Also OLS objective is easily differentiable but this is **not** the case in general!

Gradient Descent

- ▶ In class we derived:

$$\frac{\partial J(\alpha, \beta)}{\partial \alpha} = \frac{1}{N} \sum_{i=1}^N (y_i - \alpha - \beta x_i)$$

$$\frac{\partial J(\alpha, \beta)}{\partial \beta} = \frac{1}{N} \sum_{i=1}^N (y_i - \alpha - \beta x_i) x_i$$

- ▶ GD steps:

Or more concisely

$$\alpha \leftarrow \alpha - R \frac{\partial J(\alpha, \beta)}{\partial \alpha}$$

$$\beta \leftarrow \beta - R \frac{\partial J(\alpha, \beta)}{\partial \beta}$$

$$\tilde{\beta} \leftarrow \tilde{\beta} - R \nabla J(\tilde{\beta}),$$

$$\text{where } \tilde{\beta} = \begin{bmatrix} \alpha \\ \beta \end{bmatrix} \text{ and } \nabla J(\tilde{\beta}) = \begin{bmatrix} \frac{\partial J(\alpha, \beta)}{\partial \alpha} \\ \frac{\partial J(\alpha, \beta)}{\partial \beta} \end{bmatrix}$$

Gradient Descent Algorithm

Algorithm 1 GD

1: Initialize α, β , select R

2: **for** $i = 1, 2, \dots$ **do**

3:

$$\alpha \leftarrow \alpha - R \frac{\partial J(\alpha, \beta)}{\partial \alpha}$$

$$\beta \leftarrow \beta - R \frac{\partial J(\alpha, \beta)}{\partial \beta}$$

4: If convergence criterion achieved **break**

5: **end for**

Gradient Descent

- ▶ Convergence? YES! Generally in local minima
- ▶ Convergence criterion: When parameters no longer change, i.e. gradients are zero
- ▶ Choose learning rate? Small enough so that it does not diverge but large enough so that it converges fast e.g. $10^{-2}, 10^{-3} \dots$
- ▶ Fancier rules e.g. line search $R = \operatorname{argmin}_r J(x - r \nabla_x J(x))$ for $t \in \mathbb{R}_+$

Mini-Batch and Stochastic Gradient Descent

- ▶ Stochastic
- ▶ Instead of using all data points for gradient calculation use only one point
- ▶ Choose it randomly
- ▶ Diminishing step size e.g. $1/\#\text{iterations}$

Mini-Batch and Stochastic Gradient Descent

- ▶ Stochastic
 - ▶ Instead of using all data points for gradient calculation use only one point
 - ▶ Choose it randomly
 - ▶ Diminishing step size e.g. $1/\#\text{iterations}$
 - Pros: Computationally more efficient
 - Cons: Slower to converge
- ▶ Mini-Batch GD best of both worlds
 - Same notion as SGB only instead of 1 point we choose a batch K
 - typical values for K are 8, 16, 24, 32 etc

SGD and Mini-Batch GD Algorithms

Algorithm 2 Mini-Batch GD

```
1: Initialize  $\alpha, \beta$ , select  $R$ , batch-size  $K$ 
2: for  $i = 1, 2, \dots$  do
3:   Randomly shuffle data
4:   for  $j = 1, 2, \dots, \lfloor N/K \rfloor$  do
5:     Choose  $j$ th batch ( $\text{batch}(j)$ ) from data
6:
```

$$\alpha \leftarrow \alpha - R \frac{\partial J^j(\alpha, \beta)}{\partial \alpha}$$
$$\beta \leftarrow \beta - R \frac{\partial J^j(\alpha, \beta)}{\partial \beta}$$

```
7:   end for
8:   If convergence criterion achieved break
9: end for
```

where for our bivariate regression problem

$$\frac{\partial J^j(\alpha, \beta)}{\partial \alpha} = \frac{1}{K} \sum_{k \in \text{batch}(j)} y_k - \alpha - \beta x_k, \quad \frac{\partial J^j(\alpha, \beta)}{\partial \beta} = \frac{1}{K} \sum_{k \in \text{batch}(j)} (y_k - \alpha - \beta x_k) x_k \quad 10$$

Implementation Details

- ▶ Remember when shuffling your data to shuffle y and x simultaneously
- ▶ i.e. If initially

after shuffling and not

$$y = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix}, x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} \quad y = \begin{bmatrix} y_2 \\ y_3 \\ y_1 \end{bmatrix} x = \begin{bmatrix} x_2 \\ x_3 \\ x_1 \end{bmatrix} \quad y = \begin{bmatrix} y_2 \\ y_1 \\ y_3 \end{bmatrix} x = \begin{bmatrix} x_1 \\ x_3 \\ x_2 \end{bmatrix}$$

- ▶ In non-convex scenarios you should try different initialization points and compare

Linear Regression (Closed form Solution)

- ▶ Remember $\|z\|_2^2 = z^T z = z_1^2 + z_2^2 + \cdots + z_n^2 = \sum_{i=1}^N z_i^2$
- ▶ Vector calculus:
- ▶ Gradient is

$$\nabla_z J(z) = \nabla_{z_1, \dots, z_n} J(z_1, \dots, z_n) = \begin{bmatrix} \frac{\partial J(z_1, \dots, z_n)}{\partial z_1} \\ \vdots \\ \frac{\partial J(z_1, \dots, z_n)}{\partial z_n} \end{bmatrix}$$

- ▶ Also $\nabla_z z^T A z = 2A z$ (for any square symmetric matrix)
- ▶ $\nabla_z a^T z = a$ and $\nabla_z z^t a = a$ for any vector $a \in \mathbb{R}^n$

- ▶ Objective can be written as

$$J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - a - \beta x_i)^2$$

- ▶ Objective can be written as

$$J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - a - \beta x_i)^2$$

$$\frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \alpha - \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \beta \right\|_2^2$$

- ▶ Objective can be written as

$$J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - a - \beta x_i)^2$$

$$\frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \alpha - \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \beta \right\|_2^2 = \frac{1}{2N} \left\| \underbrace{\begin{bmatrix} Y \\ y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}}_{Y} - \underbrace{\begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}}_X \underbrace{\begin{bmatrix} \tilde{\alpha} \\ \tilde{\beta} \end{bmatrix}}_{\tilde{\beta}} \right\|_2^2$$

- ▶ Objective can be written as

$$J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - a - \beta x_i)^2$$

$$\begin{aligned} \frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \alpha - \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \beta \right\|_2^2 &= \frac{1}{2N} \left\| \underbrace{\begin{bmatrix} Y \\ y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}}_{Y} - \underbrace{\begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}}_X \tilde{\beta} \right\|_2^2 \\ &= \frac{1}{2N} \| Y - X \tilde{\beta} \|_2^2 \end{aligned}$$

- Objective can be written as

$$J(\alpha, \beta) = \frac{1}{2N} \sum_{i=1}^N (y_i - a - \beta x_i)^2$$

$$\frac{1}{2N} \left\| \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} - \begin{bmatrix} 1 \\ 1 \\ \vdots \\ 1 \end{bmatrix} \alpha - \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} \beta \right\|_2^2 = \frac{1}{2N} \left\| \underbrace{\begin{bmatrix} Y \\ y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}}_{Y} - \underbrace{\begin{bmatrix} 1 & x_1 \\ 1 & x_2 \\ \vdots & \vdots \\ 1 & x_n \end{bmatrix}}_X \underbrace{\begin{bmatrix} \tilde{\beta} \\ \alpha \end{bmatrix}}_{\tilde{\beta}} \right\|_2^2$$

$$= \frac{1}{2N} \|Y - X\tilde{\beta}\|_2^2 = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta})$$

Linear Regression (Closed Form Solution)

- ▶ gradient

$$\begin{aligned} J(\alpha, \beta) &= J(\tilde{\beta}) = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta}) \\ &= \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \end{aligned}$$

Linear Regression (Closed Form Solution)

- ▶ gradient

$$\begin{aligned} J(\alpha, \beta) &= J(\tilde{\beta}) = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta}) \\ &= \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \end{aligned}$$

$$\begin{aligned} \nabla_{\tilde{\beta}} J(\tilde{\beta}) &= \nabla_{\tilde{\beta}} \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \\ &= \frac{1}{2N} (2X^T X \tilde{\beta} - X^T Y - X^T Y) = \frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) \end{aligned}$$

Linear Regression (Closed Form Solution)

- ▶ gradient

$$\begin{aligned} J(\alpha, \beta) &= J(\tilde{\beta}) = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta}) \\ &= \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \end{aligned}$$

$$\begin{aligned} \nabla_{\tilde{\beta}} J(\tilde{\beta}) &= \nabla_{\tilde{\beta}} \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \\ &= \frac{1}{2N} (2X^T X \tilde{\beta} - X^T Y - X^T Y) = \frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) \end{aligned}$$

- ▶ Set derivative (gradient) equal to zero

Linear Regression (Closed Form Solution)

- ▶ gradient

$$\begin{aligned} J(\alpha, \beta) &= J(\tilde{\beta}) = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta}) \\ &= \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \end{aligned}$$

$$\begin{aligned} \nabla_{\tilde{\beta}} J(\tilde{\beta}) &= \nabla_{\tilde{\beta}} \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \\ &= \frac{1}{2N} (2X^T X \tilde{\beta} - X^T Y - X^T Y) = \frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) \end{aligned}$$

- ▶ Set derivative (gradient) equal to zero
- ▶ $\frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) = 0 \rightarrow X^T X \tilde{\beta} = X^T Y$

Linear Regression (Closed Form Solution)

- ▶ gradient

$$\begin{aligned} J(\alpha, \beta) &= J(\tilde{\beta}) = \frac{1}{2N} (Y - X\tilde{\beta})^T (Y - X\tilde{\beta}) \\ &= \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \end{aligned}$$

$$\begin{aligned} \nabla_{\tilde{\beta}} J(\tilde{\beta}) &= \nabla_{\tilde{\beta}} \frac{1}{2N} (Y^T Y + \tilde{\beta}^T X^T X \tilde{\beta} - Y^T X \tilde{\beta} - \tilde{\beta}^T X^T Y) \\ &= \frac{1}{2N} (2X^T X \tilde{\beta} - X^T Y - X^T Y) = \frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) \end{aligned}$$

- ▶ Set derivative (gradient) equal to zero
- ▶ $\frac{1}{2N} (2X^T X \tilde{\beta} - 2X^T Y) = 0 \rightarrow X^T X \tilde{\beta} = X^T Y$
- ▶ $\tilde{\beta} = (X^T X)^{-1} X^T Y$